

Chapter 9

Managing Linux Processes

Chapter Objectives

- 1 Categorize the different types of processes on a Linux system.
- 2 View processes using standard Linux utilities.
- 3 Explain the difference between common kill signals.
- 4 Describe how binary programs and shell scripts are executed.
- 5 Create and manipulate background processes.
- 6 Use standard Linux utilities to modify the priority of a process.
- 7 Schedule commands to execute in the future using the at daemon.
- 8 Schedule commands to execute repetitively using the cron daemon.


A typical Linux system can run thousands of processes simultaneously, including those that you have explored in previous chapters. In this chapter, you focus on viewing and managing processes. In the first part of the chapter, you examine the different types of processes on a Linux system and how to view them and terminate them. You then discover how processes are executed on a system, run in the background, and prioritized. Finally, you examine the various methods used to schedule commands to execute in the future.

Linux Processes

Throughout this book, the terms “program” and “process” are used interchangeably. The same is true in the workplace. However, a fine distinction exists between these two terms. Technically, a **program** is an executable file on the filesystem that can be run when you execute it. A **process**, on the other hand, is a program that is running in memory and on the CPU. In other words, a process is a program in action.

If you start a process while logged in to a terminal, that process runs in that terminal and is labeled a **user process**. Examples of user processes include `ls`, `grep`, and `find`, not to mention most of the other commands that you have executed throughout this book. Recall that a system process that is not associated with a terminal is called a **daemon process**; these processes are typically started on system startup, but you can also start them manually. Most daemon processes provide system services, such as printing, scheduling, and system maintenance, as well as network server services, such as web servers, database servers, file servers, and print servers.

Every process has a unique **process ID (PID)** that allows the kernel to identify it uniquely. In addition, each process can start an unlimited number of other processes called **child processes**. Conversely, each process must have been started by an existing process called a **parent process**. As a result, each process has a **parent process ID (PPID)**, which identifies the process that started it. An example of the relationship between parent and child processes is depicted in Figure 9-1.

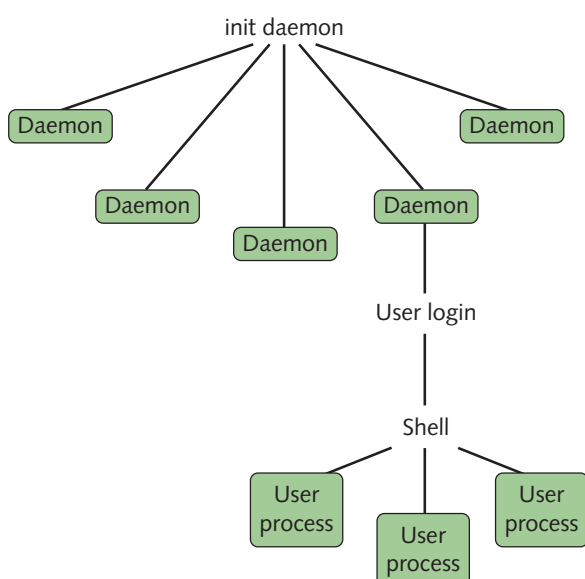
Figure 9-1 Parent and child processes**Note 1**

PIDs are not necessarily given to new processes in sequential order; each PID is generated from free entries in a process table used by the Linux kernel.

Note 2

Remember that although each process can have an unlimited number of child processes, it can only have one parent process.

The first process started by the Linux kernel is the initialize (or init) daemon, which has a PID of 1 and a PPID of 0, the latter of which refers to the kernel itself. The init daemon then starts most other daemons during the system initialization process, including those that allow for user logins. After you log in to the system, the login program starts a shell. The shell then interprets user commands and starts all user processes. Thus, each process on the Linux system can be traced back to the init daemon by examining the series of PPIDs, as shown in Figure 9-2.

Figure 9-2 Process genealogy

Note 3

The init daemon is often referred to as the “grandfather of all user processes.”

Note 4

On Linux systems that use the UNIX SysV system initialization process, the init daemon will be listed as `init` within command output. On Linux systems that use the Systemd system initialization process, the init daemon will either be listed as `init` or `systemd` within command output.

Viewing Processes

Although several Linux utilities can view processes, the most versatile and common is the **ps command**. Without arguments, the `ps` command simply displays a list of processes that are running in the current shell. The following example shows the output of this command while the user is logged in to `tty2`:

```
[root@server1 ~]# ps
  PID TTY          TIME CMD
 2159 tty2        00:00:00 bash
 2233 tty2        00:00:00 ps
[root@server1 ~]#_
```

The preceding output shows that two processes were running in the terminal `tty2` when the `ps` command executed. The command that started each process (CMD) is listed next to the time it has taken on the CPU (TIME), its PID, and terminal (TTY). In this case, the process took less than one second to run, and so the time elapsed reads nothing. To find out more about these processes, you could instead use the `-f`, or full, option to the `ps` command, as shown next:

```
[root@server1 ~]# ps -f
UID          PID  PPID  C  STIME TTY          TIME CMD
root         2159   2156  0  16:18 tty2        00:00:00 -bash
root         2233   2159  3  16:28 tty2        00:00:00 ps -f
[root@server1 ~]#_
```

This listing provides more information about each process. It displays the user who started the process (UID), the PPID, the time it was started (STIME), as well as the CPU utilization (C), which starts at zero and is incremented with each processor cycle that the process runs on the CPU.

The most valuable information provided by the `ps -f` command is each process's PPID and lineage. The `bash` process (PID = 2159) displays a shell prompt and interprets user input; it started the `ps` process (PID = 2233) because the `ps` process had a PPID of 2159.

Because daemon processes are not associated with a terminal, they are not displayed by the `ps -f` command. To display an entire list of processes across all terminals and including daemons, you can add the `-e` option to any `ps` command, as shown in the following output:

```
[root@server1 ~]# ps -ef
UID          PID  PPID  C  STIME TTY          TIME CMD
root          1      0  0  21:22 ?        00:00:00 /usr/lib/systemd/systemd
root          2      0  0  21:22 ?        00:00:00 [kthreadd]
root          3      2  0  21:22 ?        00:00:00 [ksoftirqd/0]
root          5      2  0  21:22 ?        00:00:00 [kworker/0:0H]
root          6      2  0  21:22 ?        00:00:00 [kworker/u128:0]
root          7      2  0  21:22 ?        00:00:00 [migration/0]
root          8      2  0  21:22 ?        00:00:00 [rcu_bh]
```

```

root      9      2    0 21:22 ?      00:00:00 [rcu_sched]
root     10      2    0 21:22 ?      00:00:00 [watchdog/0]
root     11      2    0 21:22 ?      00:00:00 [khelper]
root     12      2    0 21:22 ?      00:00:00 [kdevtmpfs]
root     13      2    0 21:22 ?      00:00:00 [netns]
root     14      2    0 21:22 ?      00:00:00 [writeback]
root    394      1    0 21:22 ?      00:00:00 /sbin/auditd -n
avahi   422      1    0 21:22 ?      00:00:00 avahi-daemon: running
dbus    424      1    0 21:22 ?      00:00:00 /bin/dbus-daemon --system
chrony  430      1    0 21:22 ?      00:00:00 /usr/sbin/chronyd -u
root    431      1    0 21:22 ?      00:00:00 /usr/sbin/crond -n
root    432      1    0 21:22 ?      00:00:00 /usr/sbin/atd -f
root    435      1    0 21:22 ?      00:00:00 /usr/sbin/abrttd -d -s
root    437      1    0 21:22 ?      00:00:00 /usr/bin/abrt-watch-log
root    441      1    0 21:22 ?      00:00:00 /usr/sbin/gdm
root    446      1    0 21:22 ?      00:00:00 /usr/sbin/mcelog
root    481    441    0 21:22 ?      00:00:00 /usr/libexec/gdm-simple
polkitd 482      1    0 21:22 ?      00:00:00 /usr/lib/polkit-1/polkitd
root    488    481    0 21:22 tty1    00:00:00 /usr/bin/Xorg :0
root    551    481    0 21:22 ?      00:00:00 gdm-session-worker
root    552      1    0 21:22 ?      00:00:00 /usr/sbin/NetworkManager
gdm     852    551    0 21:23 ?      00:00:00 /usr/bin/gnome-session
gdm     856      1    0 21:23 ?      00:00:00 /usr/bin/dbus-launch
gdm    1018      1    0 21:23 ?      00:00:00 /bin/dbus-daemon --fork
root   1020    552    0 21:23 ?      00:00:00 /sbin/dhclient -d -sf
gdm    1045   1018    0 21:23 ?      00:00:00 /bin/dbus-daemon
root   1072      1    0 21:23 ?      00:00:00 /usr/libexec/upowerd
gdm    1077    852    0 21:23 ?      00:00:03 gnome-shell --mode=gdm
gdm    1087      1    0 21:23 ?      00:00:00 /usr/bin/pulseaudio
gdm    1148      1    0 21:23 ?      00:00:00 /usr/libexec/goa-daemon
root   1164      1    0 21:23 ?      00:00:00 login -- root
root   1175   1164    0 21:23 tty2    00:00:00 -bash
root   1742   1175    0 21:33 tty2    00:00:00 ps -ef
[root@server1 ~]#_

```

As shown in the preceding output, the kernel thread daemon (kthreadd) has a PID of 2 and starts most subprocesses within the actual Linux kernel because those subprocesses have a PPID of 2, whereas the init daemon (/usr/lib/systemd/systemd, PID=1) starts most other daemons because those daemons have a PPID of 1. In addition, there is a ? in the TTY column for daemons and kernel subprocesses because they do not run on a terminal.

Because the output of the `ps -ef` command can be several hundred lines long on a Linux server, you usually pipe its output to the `less` command to send the output to the terminal screen page-by-page, or to the `grep` command, which can be used to display lines containing only certain information. For example, to display only the BASH shells on the system, you could use the following command:

```

[root@server1 ~]# ps -ef | grep bash
user1    2094    2008    0 14:29 pts/1    00:00:00 -bash
root     2159    2156    0 14:30 tty2      00:00:00 -bash
root     2294    2159    0 14:44 tty2      00:00:00 grep --color=auto bash
[root@server1 ~]#_

```

Notice that the `grep bash` command is also displayed alongside the BASH shells in the preceding output because it was running in memory at the time the `ps` command was executed. This might not always be the case because the Linux kernel schedules commands to run based on a variety of factors.

The `-e` and `-f` options are the most common options used with the `ps` command; however, many other options are available. The `-l` option to the `ps` command lists even more information about each process than the `-f` option. An example of using this option to view the processes in the terminal `tty2` is shown in the following output:

```
[root@server1 ~]# ps -l
F S  UID  PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
4 S   0  2159  2156  0  80   0 -  1238 wait  tty2      00:00:00 bash
4 R   0  2295  2159  2  80   0 -   744 -    tty2      00:00:00 ps
[root@server1 ~]#_
```

The process flag (F) indicates particular features of the process; the flag of 4 in the preceding output indicates that the root user ran the process. The **process state** (S) column is the most valuable to systems administrators because it indicates what the process is currently doing. If a process is not being run on the processor at the current time, you see an S (interruptible sleep) in the process state column; processes are in this state most of the time and are awoken (interrupted) by other processes when they are needed, as seen with `bash` in the preceding output. You will see an R in this column if the process is currently running on the processor, a D (uninterruptible sleep) if it is waiting for disk access, or a T if it has stopped or is being traced by another process. In addition to these, you might also see a Z in this column, indicating a **zombie process**. When a process finishes executing, the parent process must check to see if it executed successfully and then release the child process's PID so that it can be used again. While a process is waiting for its parent process to release the PID, the process is said to be in a zombie state, because it has finished but still retains a PID. On a busy Linux server, zombie processes can accumulate and prevent new processes from being created; if this occurs, you can kill the parent process of the zombies, as discussed in the next section.

Note 5

Zombie processes are also known as defunct processes.

Note 6

To view a list of zombie processes on your entire system, you could use the `ps -el | grep Z` command.

Process priority (PRI) is the priority used by the kernel for the process; it is measured between 0 (high priority) and 127 (low priority). The **nice value** (NI) can be used to affect the process priority indirectly; it is measured between -20 (a greater chance of a high priority) and 19 (a greater chance of a lower priority). The ADDR in the preceding output indicates the memory address of the process, whereas the WCHAN indicates what the process is waiting for while sleeping. In addition, the size of the process in memory (SZ) is listed and measured in kilobytes; often, it is roughly equivalent to the size of the executable file on the filesystem.

Some options to the `ps` command are not prefixed by a dash character; these are referred to as Berkeley style options. The two most common of these are the `a` option, which lists all processes across terminals, and the `x` option, which lists processes that do not run on a terminal, as shown in the following output for the first 10 processes on the system:

```
[root@server1 ~]# ps ax | head -11
PID TTY          STAT     TIME COMMAND
  1 ?            S        0:01 /usr/lib/systemd/systemd
  2 ?            S        0:00 [kthreadd]
  3 ?            S        0:00 [migration/0]
  4 ?            S<       0:00 [ksoftirqd/0]
```

```

5 ?      S      0:00 [watchdog/0]
6 ?      S      0:00 [migration/1]
7 ?      S      0:00 [ksoftirqd/1]
8 ?      S      0:00 [watchdog/1]
9 ?      S      0:00 [events/0]
10 ?     S      0:00 [events/1]
[root@server1 ~]#_

```

The columns just listed are equivalent to those discussed earlier; however, the process state column is identified with STAT and might contain additional characters to indicate the full nature of the process state. For example, a W indicates that the process has no contents in memory, a < symbol indicates a high-priority process, and an N indicates a low-priority process.

Note 7

For a full list of symbols that may be displayed in the STAT or S columns shown in prior output, consult the manual page for the `ps` command.

Several dozen options to the `ps` command can be used to display processes and their attributes; the options listed in this section are the most common and are summarized in Table 9-1.

Table 9-1 Common options to the `ps` command

Option	Description
-e	Displays all processes running on terminals as well as processes that do not run on a terminal (daemons)
-f	Displays a full list of information about each process, including the UID, PID, PPID, CPU utilization, start time, terminal, processor time, and command name
-l	Displays a long list of information about each process, including the flag, state, UID, PID, PPID, CPU utilization, priority, nice value, address, size, WCHAN, terminal, and command name
-Z	Displays SELinux context information about each process (discussed further in Chapter 14)
a	Displays all processes running on terminals
x	Displays all processes that do not run on terminals

The `ps` command is not the only command that can view process information. The kernel exports all process information subdirectories under the `/proc` directory. Each subdirectory is named for the PID of the process that it contains information for, as shown in the following output:

```

[root@server1 ~]# ls /proc
1      1174  1746  28   407  473  852      irq          slabinfo
10     1175  175   292  409  48   856      kallsyms     softirqs
1018   12    1754  3    411  481  9        kcore        stat
1020   1213  176   307  412  482  acpi     keys         swaps
1025   1216  1760  328  414  488  buddyinfo key-users    sys
1045   1220  177   350  415  49    bus      kmsg         sysrq
1052   13    178   351  418  5     cgroups  kpagecount  sysvipc
1065   14    179   353  420  551  cmdline kpageflags  timer_list
1072   15    18    354  421  552  consoles loadavg      timer_stats
1077   16    180   357  422  58    cpuinfo  locks        tty
1080   164   181   358  424  59    crypto   mdstat       uptime

```

```

1087 165 1810 359 430 6 devices meminfo version
1097 167 188 370 431 60 diskstats misc vmallocinfo
11 168 189 371 432 62 dma modules vmstat
1111 169 19 372 435 64 driver mounts zoneinfo
1114 17 191 383 437 7 execdomains mtrr
1117 170 2 384 440 72 fb net
1122 171 20 385 441 748 filesystems pagetypeinfo
1144 172 204 386 446 8 fs partitions
1148 173 205 387 45 814 interrupts sched_debug
1164 174 206 388 46 823 iomem scsi
1171 1745 276 394 47 842 ioports self
[root@server1 ~]#_

```

Thus, any program that can read from the /proc directory can display process information. For example, the **ps** command displays the lineage of a process by tracing its PPIDs until the init daemon. The first 26 lines of this command are shown in the following output:

```

[root@server1 ~]# ps -ef | head -26
systemd--ModemManager--2*[{ModemManager}]
--NetworkManager--dhclient
--3*[{NetworkManager}]
--2*[abrt-watch-log]
--abrt-d
--accounts-daemon--2*[{accounts-daemon}]
--alsactl
--at-spi-bus-laun--dbus-daemon--{dbus-daemon}
--3*[{at-spi-bus-laun}]
--at-spi2-registr--{at-spi2-registr}
--atd
--auditd--audispd--sedispatch
--{audispd}
--{auditd}
--avahi-daemon--avahi-daemon
--bluetoothd
--chronyd
--colord--2*[{colord}]
--crond
--2*[dbus-daemon--{dbus-daemon}]
--dbus-launch
--dconf-service--2*[{dconf-service}]
--firewalld--{firewalld}
--gdm--gdm-simple-slav--Xorg
--gdm-session--gnome-session--gnome-settings
--gnome-shell
[root@server1 ~]#_

```

The most common program used to display processes, aside from **ps**, is the **top** command. The **top** command displays an interactive screen listing processes organized by processor time. Processes that use the most processor time are listed at the top of the screen. An example of the screen that appears when you type the **top** command is shown next:

```

top - 21:55:15 up 32 min, 3 users, load average: 0.15, 0.06, 0.02
Tasks: 134 total, 1 running, 133 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.4 us, 0.4 sy, 0.0 ni, 95.9 id, 2.0 wa, 1.1 hi, 0.1 si, 0.0 st
MiB Mem: 7944.0 total, 7067.5 free, 467.3 used, 409.2 buff/cache
MiB Swap: 7944.0 total, 7944.0 free, 0.0 used, 7226.7 avail Mem

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1077	gdm	20	0	1518192	150684	36660	R	19.4	7.4	0:33.93	gnome-shell
2130	root	20	0	123636	1644	1180	R	0.7	0.1	0:00.05	top
1	root	20	0	51544	7348	2476	S	0.0	0.4	0:00.98	systemd


```

 2 root    20    0          0          0          0 S   0.0   0.0    0:00.00 kthreadd
 3 root    20    0          0          0          0 S   0.0   0.0    0:00.01 ksoftirqd/0
 5 root     0 -20          0          0          0 S   0.0   0.0    0:00.00 kworker/0:0H
 6 root    20    0          0          0          0 S   0.0   0.0    0:00.02 kworker/u12+
 7 root    rt     0          0          0          0 S   0.0   0.0    0:00.00 migration/0
 8 root    20    0          0          0          0 S   0.0   0.0    0:00.00 rcu_bh
 9 root    20    0          0          0          0 S   0.0   0.0    0:00.34 rcu_sched
10 root    rt     0          0          0          0 S   0.0   0.0    0:00.00 watchdog/0
11 root     0 -20          0          0          0 S   0.0   0.0    0:00.00 khelper
12 root    20    0          0          0          0 S   0.0   0.0    0:00.00 kdevtmpfs
13 root     0 -20          0          0          0 S   0.0   0.0    0:00.00 netns
14 root     0 -20          0          0          0 S   0.0   0.0    0:00.00 writeback
15 root     0 -20          0          0          0 S   0.0   0.0    0:00.00 kintegrityd
16 root     0 -20          0          0          0 S   0.0   0.0    0:00.00 bioset

```

Note that the `top` command displays many of the same columns that the `ps` command does, yet it contains a summary paragraph at the top of the screen and a cursor between the summary paragraph and the process list. From the preceding output, you can see that the `gnome-shell` uses the most processor time, followed by the `top` command itself (`top`) and the `init` daemon (`systemd`).

You might come across a process that has encountered an error during execution and continuously uses up system resources. These processes are referred to as **rogue processes** and appear at the top of the listing produced by the `top` command. The `top` command can also be used to change the priority of processes or kill them. Thus, you can stop rogue processes from the `top` command immediately after they are identified. Process priority and killing processes are discussed later in this chapter. To get a full listing of the different commands that you can use while in the `top` utility, press `h` to get a help screen.

Note 8

Rogue processes are also known as runaway processes.

Note 9

Many Linux administrators choose to install and use the **htop command** instead of `top`. The `htop` command provides the same functionality as `top` but displays results in color and includes a resource utilization graph at the top of the screen, as well as a usage legend at the bottom of the screen.

Killing Processes

As indicated earlier, a large number of rogue and zombie processes use up system resources. When system performance suffers due to these processes, you should send them a **kill signal**, which terminates a process. The most common command used to send kill signals is the **kill command**. All told, the `kill` command can send many different kill signals to a process. Each of these kill signals operates in a different manner. To view the kill signal names and associated numbers, you can use the `-l` option to the `kill` command, as shown in the following output:

```

[root@server1 ~]# kill -l
 1) SIGHUP      2) SIGINT      3) SIGQUIT     4) SIGILL
 5) SIGTRAP     6) SIGABRT    7) SIGBUS      8) SIGFPE
 9) SIGKILL    10) SIGUSR1   11) SIGSEGV    12) SIGUSR2
13) SIGPIPE    14) SIGALRM   15) SIGTERM    17) SIGCHLD
18) SIGCONT    19) SIGSTOP   20) SIGTSTP    21) SIGTTIN

```



```

22) SIGTTOU      23) SIGURG      24) SIGXCPU      25) SIGXFSZ
26) SIGVTALRM   27) SIGPROF     28) SIGWINCH     29) SIGIO
30) SIGPWR      31) SIGSYS      33) SIGRTMIN     34) SIGRTMIN+1
35) SIGRTMIN+2  36) SIGRTMIN+3  37) SIGRTMIN+4  38) SIGRTMIN+5
39) SIGRTMIN+6  40) SIGRTMIN+7  41) SIGRTMIN+8  42) SIGRTMIN+9
43) SIGRTMIN+10 44) SIGRTMIN+11 45) SIGRTMIN+12 46) SIGRTMIN+13
47) SIGRTMIN+14 48) SIGRTMIN+15 49) SIGRTMAX-15 50) SIGRTMAX-14
51) SIGRTMAX-13 52) SIGRTMAX-12 53) SIGRTMAX-11 54) SIGRTMAX-10
55) SIGRTMAX-9  56) SIGRTMAX-8  57) SIGRTMAX-7  58) SIGRTMAX-6
59) SIGRTMAX-5  60) SIGRTMAX-4  61) SIGRTMAX-3  62) SIGRTMAX-2
63) SIGRTMAX-1  64) SIGRTMAX
[root@server1 ~]#_

```

Most of the kill signals listed in the preceding output are not useful for systems administrators. The five most common kill signals used for administration are listed in Table 9-2.

Table 9-2 Common administrative kill signals

Name	Number	Description
SIGHUP	1	Also known as the hang-up signal, it stops a process, then restarts it with the same PID. If you edit the configuration file used by a running daemon, that daemon might be sent a SIGHUP to restart the process; when the daemon starts again, it reads the new configuration file.
SIGINT	2	This signal sends an interrupt signal to a process. Although this signal is one of the weakest kill signals, it works most of the time. When you use the Ctrl+c key combination to kill a currently running process, a SIGINT is actually being sent to the process.
SIGQUIT	3	Also known as a core dump, the quit signal terminates a process by taking the process information in memory and saving it to a file called core on the filesystem in the current working directory. You can use the Ctrl+\ key combination to send a SIGQUIT to a process that is currently running.
SIGTERM	15	The software termination signal is the most common kill signal used by programs to kill other processes. It is the default kill signal used by the <code>kill</code> command.
SIGKILL	9	Also known as the absolute kill signal, it forces the Linux kernel to stop executing the process by sending the process's resources to a special device file called <code>/dev/null</code> .

To send a kill signal to a process, you specify the kill signal to send as an option to the `kill` command, followed by the appropriate PID of the process. For example, to send a SIGQUIT to a process called `sample`, you could use the following commands to locate and terminate the process:

```

[root@server1 ~]# ps -ef | grep sample
root      1199      1  0 Jun30 tty3      00:00:00 /sbin/sample
[root@server1 ~]# kill -3 1199
[root@server1 ~]#_
[root@server1 ~]# ps -ef | grep sample
[root@server1 ~]#_

```

Note 10

The `kill -SIGQUIT 1199` command does the same thing as the `kill -3 1199` command shown in the preceding output.

Note 11

If you do not specify the kill signal when using the `kill` command, the `kill` command uses the default kill signal, the `SIGTERM` signal.

Note 12

You can also use the **pidof command** to find the PID of a process to use as an argument to the `kill` command. For example, `pidof sample` will return the PID of the sample process.

When sending a kill signal to several processes, it is often easier to locate the process PIDs using the **pgrep command**. The `pgrep` command returns a list of PIDs for processes that match a regular expression, or other criteria. For example, to send a `SIGQUIT` to all processes started by the user `bini` whose process name starts with the letters `psql`, you could use the following commands to locate and terminate the process:

```
[root@server1 ~]# pgrep -u bini "^psql"
1344
1501
1522
[root@server1 ~]# kill -3 1344 1501 1522
[root@server1 ~]# _
[root@server1 ~]# pgrep -u bini "^psql"
[root@server1 ~]# _
```

Some processes have the ability to ignore, or **trap**, certain kill signals that are sent to them. The only kill signal that cannot be trapped by any process is the `SIGKILL`. Thus, if a `SIGINT`, `SIGQUIT`, and `SIGTERM` do not terminate a stubborn process, you can use a `SIGKILL` to terminate it. However, you should only use `SIGKILL` as a last resort because it prevents a process from closing open files and other resources properly.

Note 13

You can use the **lsof (list open files) command** to view the files that a process has open before sending a `SIGKILL`. For example, to see the files that are used by the process with PID 1399, you can use the `lsof -p 1399` command.

If you send a kill signal to a process that has children, the parent process terminates all of its child processes before terminating itself. Thus, to kill several related processes, you can simply send a kill signal to their parent process. In addition, to kill a zombie process, it is often necessary to send a kill signal to its parent process.

Note 14

To prevent a child process from being terminated when the parent process is terminated, you can start the child process with the **nohup command**. For example, executing the `nohup catheona` command within your shell would execute the `catheona` child process without any association to the parent shell process that started it.

Another command that can be used to send kill signals to processes is the **killall** command. The **killall** command works similarly to the **kill** command in that it takes the kill signal as an option; however, it uses the process name to kill instead of the PID. This allows multiple processes of the same name to be killed in one command. An example of using the **killall** command to send a SIGQUIT to multiple **sample** processes is shown in the following output:

```
[root@server1 ~]# ps -ef | grep sample
root      1729      1  0 Jun30 tty3      00:00:00 /sbin/sample
root      20198     1  0 Jun30 tty4      00:00:00 /sbin/sample
[root@server1 ~]# killall -3 sample
[root@server1 ~]# _
[root@server1 ~]# ps -ef | grep sample
[root@server1 ~]# _
```

Note 15

Alternatively, you could use the command **killall -SIGQUIT sample** to do the same as the **killall -3 sample** command used in the preceding output.

Note 16

As with the **kill** command, if you do not specify the kill signal when using the **killall** command, it sends a SIGTERM signal by default.

You can also use the **pkill** command to kill processes by process name. However, the **pkill** command allows you to identify process names using regular expressions as well as specify other criteria. For example, the **pkill -u bini -3 "^psql"** command will send a SIGQUIT signal to processes started by the **bini** user that begin with the letters **psql**.

In addition to the **kill**, **killall**, and **pkill** commands, the **top** command can be used to kill processes. While in the **top** utility, press the **k** key and supply the appropriate PID and kill signal when prompted.

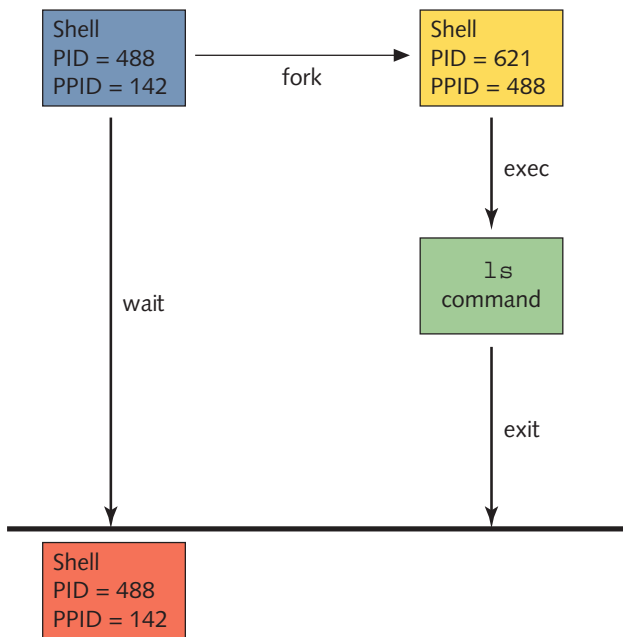
Process Execution

You can execute three main types of Linux commands:

- Binary programs
- Shell scripts
- Shell functions

Most commands, such as **ls**, **find**, and **grep**, are binary programs that exist on the filesystem until executed. They were written in a certain programming language and compiled into a binary format that only the computer can understand. Other commands, such as **cd** and **exit**, are built into the shell running in memory, and they are called shell functions. Shell scripts can also contain a list of binary programs, shell functions, and special constructs for the shell to execute in order.

When executing compiled programs or shell scripts, the shell that interprets the command you typed creates a new shell. This creation of a new subshell is known as **forking** and is carried out by the **fork** function in the shell. The new subshell then executes the binary program or shell script using its **exec** function. After the binary program or shell script has completed, the new shell uses its **exit** function to kill itself and return control to the original shell. The original shell uses its **wait** function to wait for the new shell to carry out the aforementioned tasks before returning a prompt to the user. Figure 9-3 depicts this process when a user types the **ls** command at the command line.

Figure 9-3 Process forking

Running Processes in the Background

As discussed in the previous section, the shell creates, or forks, a subshell to execute most commands on the Linux system. Unfortunately, the original shell must wait for the command in the subshell to finish before displaying a shell prompt to accept new commands. Commands run in this fashion are known as **foreground processes**.

Alternatively, you can omit the wait function shown in Figure 9-3 by appending an ampersand (&) character to the command. Commands run in this fashion are known as **background processes**. When a command is run in the background, the shell immediately returns the shell prompt for the user to enter another command. To run the `sample` command in the background, you can enter the following command:

```
[root@server1 ~]# sample &
[1] 2583
[root@server1 ~]#_
```

Note 17

Space characters between the command and the ampersand (&) are optional. In other words, the command `sample&` is equivalent to the command `sample &` used in the preceding output.

The shell returns the PID (2583 in the preceding example) and the background job ID (1 in the preceding example) so that you can manipulate the background job after it has been run. After the process has been started, you can use the `ps` command to view the PID or the **jobs** command to view the background job ID, as shown in the following output:

```
[root@server1 ~]# jobs
[1]+  Running                  sample &
[root@server1 ~]# ps | grep sample
2583 tty2    00:00:00 sample
[root@server1 ~]#_
```

To terminate the background process, you can send a kill signal to the PID (as shown earlier in this chapter), or you can send a kill signal to the background job ID. Background job IDs must be prefixed with a % character. To send the sample background process created earlier a SIGINT signal, you could use the following `kill` command:

```
[root@server1 ~]# jobs
[1]+  Running                  sample &
[root@server1 ~]# kill -2 %1
[1]+  Interrupt                sample
[root@server1 ~]# jobs
[root@server1 ~]# _
```

Note 18

You can also use the `killall -2 sample` command or the `top` utility to terminate the sample background process used in the preceding example.

After a background process has been started, you can move it to the foreground by using the **fg (foreground) command** followed by the background job ID. Similarly, you can pause a foreground process by using the `Ctrl+z` key combination. You can then send the process to the background with the **bg (background) command**. The `Ctrl+z` key combination assigns the foreground process a background job ID that is then used as an argument to the `bg` command. To start a sample process in the background and move it to the foreground, then pause it and move it to the background again, you can use the following commands:

```
[root@server1 ~]# sample &
[1] 7519
[root@server1 ~]# fg %1
sample
```

Ctrl+z

```
[1]+  Stopped                  sample
[root@server1 ~]# bg %1
[1]+ sample &
[root@server1 ~]# jobs
[1]+  Running                  sample &
[root@server1 ~]# _
```

When there are multiple background processes executing in the shell, the `jobs` command indicates the most recent one with a + symbol, and the second most recent one with a – symbol. If you place the % notation in a command without specifying the background job ID, the command operates on the most recent background process. An example of this is shown in the following output, in which four sample processes are started and sent SIGQUIT kill signals using the % notation:

```
[root@server1 ~]# sample &
[1] 7605
[root@server1 ~]# sample2 &
[2] 7613
[root@server1 ~]# sample3 &
[3] 7621
[root@server1 ~]# sample4 &
[4] 7629
[root@server1 ~]# jobs
```

```

[1] Running sample &
[2] Running sample2 &
[3]- Running sample3 &
[4]+ Running sample4 &
[root@server1 ~]# kill -3 %
[root@server1 ~]# jobs
[1] Running sample &
[2]- Running sample2 &
[3]+ Running sample3 &
[root@server1 ~]# kill -3 %
[root@server1 ~]# jobs
[1]- Running sample &
[2]+ Running sample2 &
[root@server1 ~]# kill -3 %
[root@server1 ~]# jobs
[1]+ Running sample &
[root@server1 ~]# kill -3 %
[root@server1 ~]# jobs
[root@server1 ~]# _

```

Process Priorities

Recall that Linux is a multitasking operating system. That is, it can perform several tasks at the same time. Because most computers contain only a single CPU, Linux executes small amounts of each process on the processor in series. This makes it seem to the user as if processes are executing simultaneously. The amount of time a process has to use the CPU is called a **time slice**; the more time slices a process has, the more time it has to execute on the CPU and the faster it executes. Time slices are typically measured in milliseconds. Thus, several hundred processes can be executing on the processor in a single second.

The `ps -l` command lists the Linux kernel priority (PRI) of a process. This value is directly related to the amount of time slices a process has on the CPU. A PRI of 0 is the most likely to get time slices on the CPU, and a PRI of 127 is the least likely to receive time slices on the CPU. An example of this command is shown next:

```

[root@server1 ~]# ps -l
F S  UID  PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY  TIME CMD
4 S   0   3194  3192  0   75   0 -  1238 wait4 pts/1  00:00:00 bash
4 S   0   3896  3194  0   76   0 -   953 - pts/1  00:00:00 sleep
4 S   0   3939  3194 13   75   0 -  7015 - pts/1  00:00:01 gedit
4 R   0   3940  3194  0   77   0 -   632 - pts/1  00:00:00 ps
[root@server1 ~]# _

```

The `bash`, `sleep`, `gedit`, and `ps` processes all have different PRI values because the kernel automatically assigns time slices based on several factors. You cannot change the PRI directly, but you can influence it indirectly by assigning a certain `nice` value to a process. A negative `nice` value increases the likelihood that the process will receive more time slices, whereas a positive `nice` value does the opposite. The range of `nice` values is depicted in Figure 9-4.


```

4 S      0   1990   1229   0   69   0   -   483 nanosl pts/0  00:00:00 sample
4 S      0   2180   1229   0   70   0   -   483 nanosl pts/0  00:00:00 sample
4 S      0   2181   1229   0   71   0   -   483 nanosl pts/0  00:00:00 sample
4 R      0   2196   1229   0   75   0   -   768 -      pts/0  00:00:00 ps
[root@server1 ~]#_

```

To lower priority of the first two sample processes by changing the nice value from 0 to +15 and view the new values, you can execute the following commands:

```

[root@server1 ~]# renice +15 1990 2180
1990 (process ID) old priority 0, new priority 15
2180 (process ID) old priority 0, new priority 15
[root@server1 ~]# ps -l
F S  UID      PID  PPID  C PRI  NI ADDR      SZ WCHAN  TTY          TIME CMD
4 S      0    1229    1228   0   71   0   -    617 wait4  pts/0    00:00:00 bash
4 S      0    1990    1229   0   93  15   -    483 nanosl  pts/0    00:00:00 sample
4 S      0    2180    1229   0   96  15   -    483 nanosl  pts/0    00:00:00 sample
4 S      0    2181    1229   0   71   0   -    483 nanosl  pts/0    00:00:00 sample
4 R      0    2196    1229   0   75   0   -    768 -      pts/0    00:00:00 ps
[root@server1 ~]#_

```

Note 20

You can also use the `top` utility to change the nice value of a running process. Press the `r` key, then supply the PID and the nice value when prompted.

Note 21

As with the `nice` command, only the root user can change the nice value to a negative value using the `renice` command.

The root user can use the `renice` command to change the priority of all processes that are owned by a certain user or group. To change the nice value to +15 for all processes owned by the users `mary` and `bini`, you could execute the command `renice +15 -u mary bini` at the command prompt. Similarly, to change the nice value to +15 for all processes started by members of the group `sys`, you could execute the command `renice +15 -g sys` at the command prompt.

Scheduling Commands

Although most processes are begun by users executing commands while logged in to a terminal, at times you might want to schedule a command to execute at some point in the future. For example, scheduling system maintenance commands to run during nonworking hours is good practice, as it does not disrupt normal business activities.

You can use two different daemons to schedule commands: the **at daemon (atd)** and the **cron daemon (crond)**. The at daemon can be used to schedule a command to execute once in the future, whereas the cron daemon is used to schedule a command to execute repeatedly in the future.

Scheduling Commands with atd

To schedule a command or set of commands for execution at a later time by the at daemon, you can specify the time as an argument to the **at command**; some common time formats used with the at command are listed in Table 9-3.

Table 9-3 Common `at` commands

Command	Description
<code>at 10:15pm</code>	Schedules commands to run at 10:15 PM on the current date
<code>at 10:15pm July 15</code>	Schedules commands to run at 10:15 PM on July 15
<code>at midnight</code>	Schedules commands to run at midnight on the current date
<code>at noon July 15</code>	Schedules commands to run at noon on July 15
<code>at teatime</code>	Schedules commands to run at 4:00 PM on the current date
<code>at tomorrow</code>	Schedules commands to run the next day
<code>at now + 5 minutes</code>	Schedules commands to run in five minutes
<code>at now + 10 hours</code>	Schedules commands to run in 10 hours
<code>at now + 4 days</code>	Schedules commands to run in four days
<code>at now + 2 weeks</code>	Schedules commands to run in two weeks
<code>at now</code> <code>at batch</code>	Schedules commands to run immediately
<code>at 9:00am 01/03/2023</code> <code>at 9:00am 01032023</code> <code>at 9:00am 03.01.2023</code>	Schedules commands to run at 9:00 AM on January 3, 2023

After being invoked, the `at` command displays an `at>` prompt allowing you to type commands to be executed, one per line. After the commands have been entered, use the `Ctrl+d` key combination to schedule the commands using `atd`.

Note 22

The `at` daemon uses the current shell's environment when executing scheduled commands. The shell environment and scheduled commands are stored in the `/var/spool/at` directory on Fedora systems and the `/var/spool/cron/atjobs` directory on Ubuntu systems.

Note 23

If the standard output of any command scheduled using `atd` has not been redirected to a file, it is normally mailed to the user. You can check your local mail by typing `mail` at a command prompt. Because most modern Linux distributions do not install a mail daemon by default, it is important to ensure that the output of any commands scheduled using `atd` are redirected to a file.

To schedule the commands `date` and `who` to run at 10:15 PM on July 15, you can use the following commands:

```
[root@server1 ~]# at 10:15pm July 15
at> date > /root/atfile
at> who >> /root/atfile
at> Ctrl+d
job 1 at Wed Jul 15 22:15:00 2023
[root@server1 ~]#_
```

As shown in the preceding output, the `at` command returns an `at` job ID. You can use this ID to query or remove the scheduled command. To display a list of `at` job IDs, you can specify the `-l` option to the `at` command:

```
[root@server1 ~]# at -l
1          Wed Jul 15 22:15:00 2023 a root
[root@server1 ~]#_
```

Note 24

Alternatively, you can use the **atq** command to see scheduled `at` jobs. The `atq` command is simply a shortcut to the `at -l` command.

Note 25

When running the `at -l` command, a regular user only sees their own scheduled `at` jobs; however, the root user sees all scheduled `at` jobs.

To see the contents of the `at` job listed in the previous output alongside the shell environment at the time the `at` job was scheduled, you can use the `-c` option to the `at` command and specify the appropriate `at` job ID:

```
[root@server1 ~]# at -c 1
#!/bin/sh
# atrun uid=0 gid=0
# mail root 0
umask 22
XDG_VTNR=2; export XDG_VTNR
XDG_SESSION_ID=1; export XDG_SESSION_ID
HOSTNAME=server1; export HOSTNAME
SHELL=/bin/bash; export SHELL
HISTSIZE=1000; export HISTSIZE
QT_GRAPHICSSYSTEM_CHECKED=1; export QT_GRAPHICSSYSTEM_CHECKED
USER=root; export USER
MAIL=/var/spool/mail/root; export MAIL
PATH=/usr/local/sbin:/usr/local/bin:/sbin:/bin:/usr/sbin:/usr/bin:/root
/bin; export PATH
PWD=/root; export PWD
LANG=en_US.UTF-8; export LANG
KDEDIRS=/usr; export KDEDIRS
HISTCONTROL=ignoredups; export HISTCONTROL
SHLVL=1; export SHLVL
XDG_SEAT=seat0; export XDG_SEAT
HOME=/root; export HOME
LOGNAME=root; export LOGNAME
LESSOPEN=\\|\\|/usr/bin/lesspipe.sh\ %s; export LESSOPEN
XDG_RUNTIME_DIR=/run/user/0; export XDG_RUNTIME_DIR
cd /root || {
    echo 'Execution directory inaccessible' >&2
    exit 1
}
${SHELL:-/bin/sh} << 'marcinDELIMITER2b2a920e'
date >/root/atfile
who >>/root/atfile
```

```
marcinDELIMITER2b2a920e
```

```
[root@server1 ~]#_
```

To remove the at job used in the preceding example, specify the `-d` option to the `at` command, followed by the appropriate at job ID, as shown in the following output:

```
[root@server1 ~]# at -d 1
```

```
[root@server1 ~]# at -l
```

```
[root@server1 ~]#_
```

Note 26

Alternatively, you can use the `atrm 1` command to remove the first at job. The `atrm` command is simply a shortcut to the `at -d` command.

If there are many commands to be scheduled using the at daemon, you can place these commands in a shell script and then schedule the shell script to execute at a later time using the `-f` option to the at command. An example of scheduling a shell script called `myscript` using the at command is shown next:

```
[root@server1 ~]# cat myscript
```

```
#this is a sample shell script
```

```
date > /root/atfile
```

```
who >> /root/atfile
```

```
[root@server1 ~]# at 10:15pm July 16 -f myscript
```

```
job 2 at Wed Jul 15 22:15:00 2023
```

```
[root@server1 ~]#_
```

If the `/etc/at.allow` and `/etc/at.deny` files do not exist, only the root user is allowed to schedule tasks using the at daemon. To give this ability to other users, create an `/etc/at.allow` file and add the names of users allowed to use the at daemon, one per line. Conversely, you can use the `/etc/at.deny` file to deny certain users access to the at daemon; any user not listed in this file is then allowed to use the at daemon. If both files exist, the system checks the `/etc/at.allow` file and does not process the entries in the `/etc/at.deny` file.

Note 27

On Fedora systems, only an `/etc/at.deny` file exists by default. Because this file is initially left blank, all users are allowed to use the at daemon. On Ubuntu systems, only an `/etc/at.deny` file exists by default, and lists daemon user accounts. As a result, the root user and other regular user accounts are allowed to use the at daemon.

Scheduling Commands with cron

The at daemon is useful for scheduling tasks that occur on a certain date in the future but is ill suited for scheduling repetitive tasks, because each task requires its own at job ID. The cron daemon is better suited for repetitive tasks because it uses configuration files called **cron tables** to specify when a command should be executed.

A cron table includes six fields separated by space or tab characters. The first five fields specify the times to run the command, and the sixth field is the absolute pathname to the command to be executed. As with the at command, you can place commands in a shell script and schedule the shell script to run repetitively; in this case, the sixth field is the absolute pathname to the shell script. Each of the fields in a cron table is depicted in Figure 9-5.

Figure 9-5 User cron table format

1	2	3	4	5	command
---	---	---	---	---	---------

1 = Minute past the hour (0–59)
2 = Hour (0–23)
3 = Day of month (1–31)
4 = Month of year (1–12)
5 = Day of week
 0 = Sun (or 7 = Sun)
 1 = Mon
 2 = Tues
 3 = Wed
 4 = Thurs
 5 = Fri
 6 = Sat

Thus, to execute the `/root/myscript` shell script at 5:20 PM and 5:40 PM Monday to Friday regardless of the day of the month or month of the year, you could use the cron table depicted in Figure 9-6.

Figure 9-6 Sample user cron table entry

1	2	3	4	5	command
---	---	---	---	---	---------

20,40	17	*	*	1–5	/root/myscript
-------	----	---	---	-----	----------------

The first field in Figure 9-6 specifies the minute past the hour. Because the command must be run at 20 minutes and 40 minutes past the hour, this field has two values, separated by a comma. The second field specifies the time in 24-hour format, with 5 PM being the 17th hour. The third and fourth fields specify the day of month and month of year, respectively, to run the command. Because the command might run during any month regardless of the day of month, both fields use the `*` wildcard shell metacharacter to match all values. The final field indicates the day of the week to run the command; as with the first field, the command must be run on multiple days, but a range of days was specified (day 1 to day 5).

Two types of cron tables are used by the cron daemon: user cron tables and system cron tables. User cron tables represent tasks that individual users schedule and exist in the `/var/spool/cron` directory on Fedora systems and the `/var/spool/cron/crontabs` directory on Ubuntu systems. System cron tables contain system tasks and exist in the `/etc/crontab` file as well as the `/etc/cron.d` directory.

User Cron Tables

On a newly installed Fedora system, all users have the ability to schedule tasks using the cron daemon because the `/etc/cron.deny` file has no contents. However, if you create an `/etc/cron.allow` file and add a list of users to it, only those users will be able to schedule tasks using the cron daemon. All other users are denied. Conversely, you can modify the `/etc/cron.deny` file to list those users who are denied the ability to schedule tasks. Thus, any users not listed in this file are allowed to schedule tasks. If both files exist, only the `/etc/cron.allow` file is processed. If neither file exists, all users are allowed to schedule tasks, which is the case on a newly installed Ubuntu system.

To create or edit a user cron table, you can use the `-e` option to the `crontab` command, which opens the nano editor by default on Fedora systems. You can then enter the appropriate cron table entries. Suppose, for example, that the root user executed the `crontab -e` command on a Fedora system. To schedule `/bin/command1` to run at 4:30 AM every Friday and `/bin/command2` to run at 2:00 PM on the first day of every month, you can add the following lines while in the nano editor:

```
30 4 * * 5 /bin/command1
0 14 1 * * /bin/command2
```

Note 28

When you run the `crontab -e` command on an Ubuntu system, you are prompted for the editor to use.

When the user saves the changes and quits the nano editor, the information is stored in the file `/var/spool/cron/username`, where *username* is the name of the user who executed the `crontab -e` command. In the preceding example, the file would be named `/var/spool/cron/root`.

To list your user cron table, you can use the `-l` option to the `crontab` command. The following output lists the cron table created earlier:

```
[root@server1 ~]# crontab -l
30 4 * * 5 /bin/command1
0 14 1 * * /bin/command2
[root@server1 ~]#_
```

Furthermore, to remove a cron table and all scheduled jobs, you can use the `-r` option to the `crontab` command, as illustrated next:

```
[root@server1 ~]# crontab -r
[root@server1 ~]# crontab -l
no crontab for root
[root@server1 ~]#_
```

The root user can edit, list, or remove any other user's cron table by using the `-u` option to the `crontab` command followed by the user name. For example, to edit the cron table for the user *mary*, the root user could use the command `crontab -e -u mary` at the command prompt. Similarly, to list and remove *mary*'s cron table, the root user could execute the commands `crontab -l -u mary` and `crontab -r -u mary`, respectively.

System Cron Tables

Linux systems are typically scheduled to run many commands during nonbusiness hours. These commands might perform system maintenance, back up data, or run CPU-intensive programs. While Systemd timer units can be configured to run these commands, they are often scheduled by the cron daemon from entries in the system cron table `/etc/crontab`, which can only be edited by the root user. The default `/etc/crontab` file on a Fedora system is shown in the following output:

```
[root@server1 ~]# cat /etc/crontab
SHELL=/bin/bash
PATH=/sbin:/bin:/usr/sbin:/usr/bin
MAILTO=root

# For details see man 4 crontabs

# Example of job definition:
# .----- minute (0 - 59)
# | .----- hour (0 - 23)
# | | .----- day of month (1 - 31)
# | | | .----- month (1 - 12) OR jan,feb,mar,apr ...
# | | | | .---- day of week (0 - 6) (Sunday=0 or 7) OR
# | | | | | sun,mon,tue,wed,thu,fri,sat
# | | | | |
# * * * * * user-name command to be executed

[root@server1 ~]#_
```

The initial section of the cron table specifies the environment used while executing commands. The remainder of the file contains comments that identify the format of a cron table entry. If you add your own cron table entries to the bottom of this file, they will be executed as the root user. Alternatively, you may prefix the command within a system cron table entry with the user account that it should be executed as. For example, the `/bin/cleanup` command shown in the following output will be executed as the apache user every Friday at 11:30 PM:

```
[root@server1 ~]# cat /etc/crontab
SHELL=/bin/bash
PATH=/sbin:/bin:/usr/sbin:/usr/bin
MAILTO=root

# For details see man 4 crontabs

# Example of job definition:
# .----- minute (0 - 59)
# | .----- hour (0 - 23)
# | | .----- day of month (1 - 31)
# | | | .----- month (1 - 12) OR jan,feb,mar,apr ...
# | | | | .---- day of week (0 - 6) (Sunday=0 or 7) OR
# | | | | | sun,mon,tue,wed,thu,fri,sat
# | | | | |
# * * * * * user-name command to be executed

30 23 * * 5 apache /bin/cleanup
```

```
[root@server1 ~]#_
```

You can also place a cron table with the same information in the `/etc/cron.d` directory. Any cron tables found in this directory can have the same format as `/etc/crontab` and are run by the system. In the following example, the cron daemon is configured to run the `sa1` command every 10 minutes as the root user, and the `sa2` command at 11:53 PM as the root user each day:

```
[root@server1 ~]# cat /etc/cron.d/sysstat
# Run system activity accounting tool every 10 minutes
*/10 * * * * root /usr/lib/sa/sa1 -S DISK 1 1
# Generate a daily summary of process accounting at 23:53
53 23 * * * root /usr/lib/sa/sa2 -A
[root@server1 ~]#_
```

Note 29

The **watch command** can be used instead of the cron daemon for scheduling very frequent tasks. For example, to run the `sa1` command shown in the previous output every 1 minute (60 seconds), you could run the `watch -n 60 /usr/lib/sa/sa1 -S DISK 1 1` command.

Many administrative tasks are performed on an hourly, daily, weekly, or monthly basis. If you have a task of this type, you don't need to create a system cron table. Instead, you can place a shell script that runs the appropriate commands in one of the following directories:

- Scripts that should be executed hourly in the `/etc/cron.hourly` directory
- Scripts that should be executed daily in the `/etc/cron.daily` directory
- Scripts that should be executed weekly in the `/etc/cron.weekly` directory
- Scripts that should be executed monthly in the `/etc/cron.monthly` directory

On Fedora systems, the cron daemon runs the `/etc/cron.d/0hourly` script, which executes the contents of the `/etc/cron.hourly` directory 1 minute past the hour, every hour on the hour. The `/etc/cron.hourly/0anacron` file starts the **anacron daemon**, which then executes the contents of the `/etc/cron.daily`, `/etc/cron.weekly`, and `/etc/cron.monthly` directories at the times specified in `/etc/anacrontab`.

On Ubuntu systems, cron table entries within the `/etc/crontab` file are used to execute the contents of the `/etc/cron.hourly`, as well as the contents of the `/etc/cron.daily`, `/etc/cron.weekly`, and `/etc/cron.monthly` directories using the anacron daemon.

Note 30

If the computer is powered off during the time of a scheduled task, the cron daemon will simply not execute the task. This is why cron tables within the `/etc/cron.daily`, `/etc/cron.weekly`, and `/etc/cron.monthly` directories are often executed by the anacron daemon, which will resume task execution at the next available time if the computer is powered off during the time of a scheduled task.

Summary

- Processes are programs that are executing on the system.
- User processes are run in the same terminal as the user who executed them, whereas daemon processes are system processes that do not run on a terminal.
- Every process has a parent process associated with it and, optionally, several child processes.
- Process information is stored in the `/proc` filesystem. You can use the `ps`, `pstree`, `pgrep`, and `top` commands to view this information.
- Zombie and rogue processes that exist for long periods of time use up system resources and should be killed to improve system performance.
- You can send kill signals to a process using the `kill`, `killall`, `pkill`, and `top` commands.
- The shell creates, or forks, a subshell to execute most commands.
- Processes can be run in the background by appending an `&` to the command name. The shell assigns each background process a background job ID such that it can be manipulated afterward.
- The priority of a process can be affected indirectly by altering its nice value; nice values range from `-20` (high priority) to `+19` (low priority). Only the root user can increase the priority of a process.
- You can use the `at` and `cron` daemons to schedule commands to run at a later time. The `at` daemon schedules tasks to occur once at a later time, whereas the `cron` daemon uses cron tables to schedule tasks to occur repetitively in the future.

Key Terms

anacron daemon

at command

at daemon (atd)

atq command

atrm command

background process

bg (background) command

child process

cron daemon (crond)

cron table

crontab command

daemon process

fg (foreground) command

foreground process

forking

htop command

jobs command

kill command

kill signal

killall command

lsof (list open files) command

nice command

nice value

nohup command

parent process

parent process ID (PPID)

pgrep command

pidof command

pkill command

process

process ID (PID)

process priority

process state

program

ps command

pstree command

renice command

rogue process

time slice

top command

trap

user process

watch command

zombie process